

# Lab3 : a LSTM Cell for Image Captioning

Department of Computer Science, NCTU

TA Ziv(鍾嘉峻)

8	Factor Models + EM + Autoencoders (AE)	CNN + RNN: Image Captioning	Lab 3講解	V		V Lab2	鍾嘉峻	9/16	彭文孝
9	Generative Adversarial Networks (GAN)	CNN + RNN: Image Captioning	Lab 4講解	V	V		陳璽存	9/23	邱維辰
10	Generative Adversarial Networks (GAN)	VAE					無Lab	9/30	邱維辰
11	Final project proposal review	Final project proposal review					無Lab	10/7	彭文孝、吳毅成 邱維辰、陳永昇
12	Generative Adversarial Networks	VAE			V	Lab3	鍾嘉峻	10/14	邱維辰
13	Generative Adversarial Networks (GAN)	GAN (DC-GAN)			V		陳璽存	10/21	邱維辰
14	Reinforcement Learning (RL)	GAN (DC-GAN)	Lab5 講解	V		Lab4	李毅倫	10/28	吳毅成
15	Reinforcement Learning (RL)	RL			V		李懿倫	11/4	吳毅成
16	Reinforcement Learning (RL)	RL	Lab6 講解	V		Lab5	鄭余玄	11/11	吳毅成
17					V		鄭余玄		提供有問題的人 來問問題

# Useful Link

## General

---

- Ask anything you want in this repo~~

## NCTU DL Final Project Demo

---

- [2018 Spring](#)
- [2018 Summer](#)
- [2019 Spring](#)

## Useful Link

---

- [awesome-AI-books](#)
- [3D ML](#)
- [CNN training skill](#)

# Important Rules

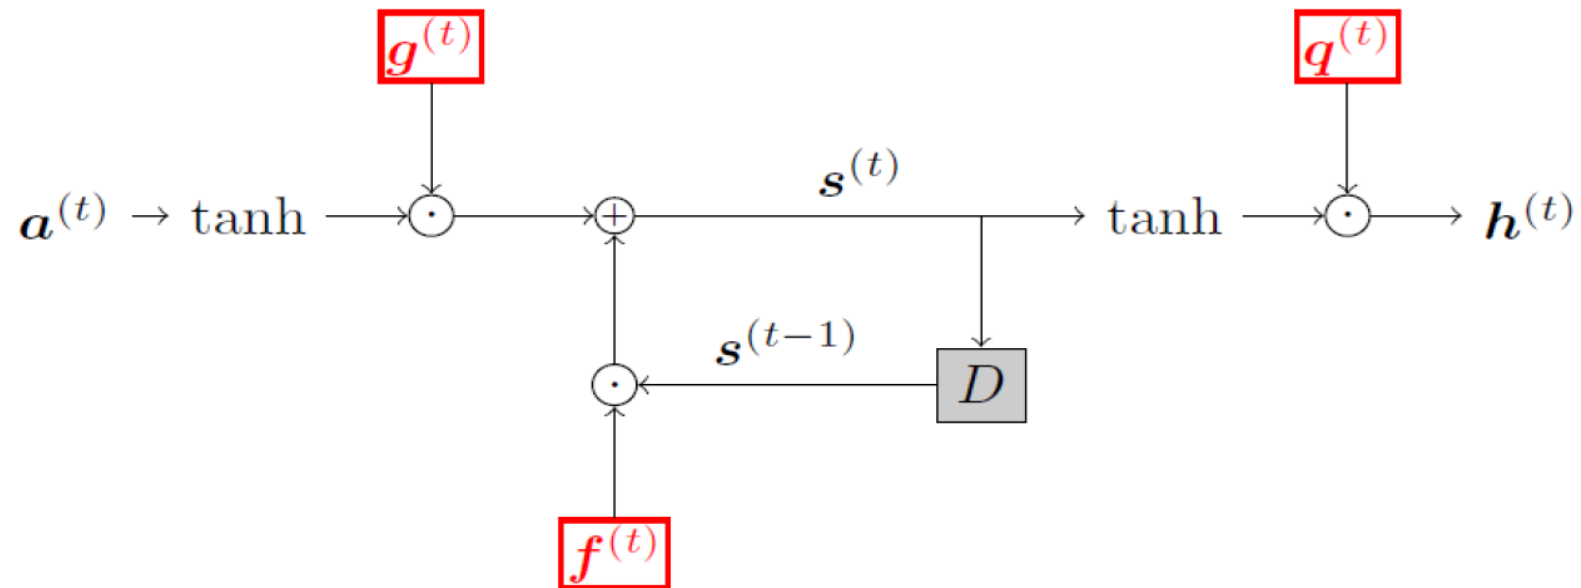
- Important Date :
  - Report Submission Deadline: 10/14 (Wed) 11:59 AM
  - Demo date: 10/14 (Wed)
- Turn in :
  - Experiment Report (.pdf)
  - Source code (.py)
- Notice: zip all files in one file and name it like 「DLP\_LAB3\_your studentID\_name.zip」, ex: 「DLP\_LAB3\_0756172\_鍾嘉峻.zip」

# Important Rules

- Email To :
  - [zhongturtle@gmail.com](mailto:zhongturtle@gmail.com)
  - Don't CC other TA
- Email Title :
  - DLP\_LAB3\_your studentID\_name
- Do not submit your weight or dataset!!
  - But you should save the model weight for demo

# Lab Objective

- In this lab, you only have to implement a LSTM cell by yourselves
  - Only `DIY_LSTM.py`
- And train an image caption model with your own LSTM cell



# Lab Requirement

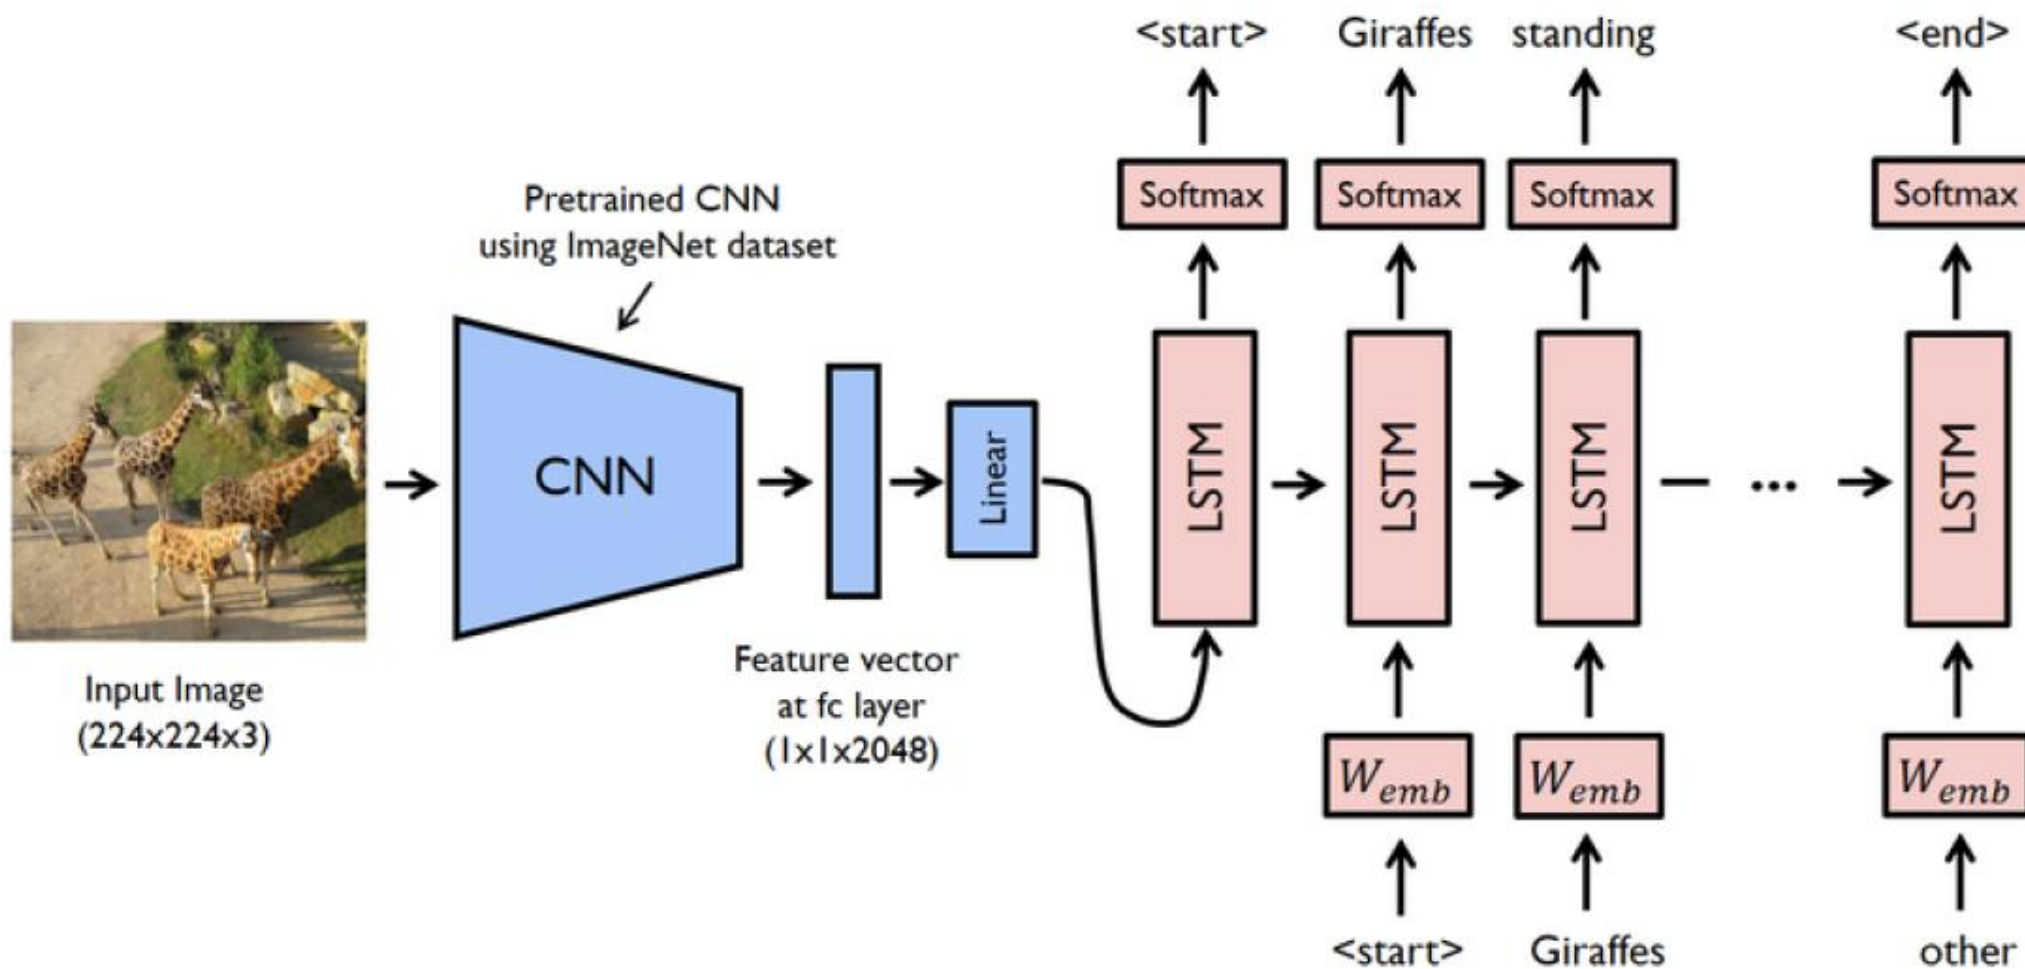
- Implement a LSTM cell
  - Please finish
  - Only Forward part , don't worried
- Replace the LSTM cell in Pytorch image-caption example with the one you implement.
- Train an image caption model

# Lab Resource & Instruction


- Clone <https://github.com/2020-DL-Training-Program/Lab3-Image-Caption.git>
  - Already on the sever
  - Please follow the Usage in Readme if you want to try on you own machine
- Get the data
  - Already on the sever
- Implement a LSTM cell (DIY\_LSTM.py)
- Train the model
- You only need to start at “3. Preprocessing” in the Readme




# Image-Caption



# Encoder-Decoder



# Encoder-Decoder



# Dataset

- ImageNet : Large Scale Visual Recognition Challenge 2012 (ILSVRC2012)
- Very good dataset for
  - Classification
  - localization
  - .....
- Dataset for this lab
  - Image
  - Annotation

## IMAGENET Large Scale Visual Recognition Challenge 2012 (ILSVRC2012)

Held in conjunction with [PASCAL Visual Object Classes Challenge 2012 \(VOC2012\)](#).

[Introduction](#) [Task](#) [Timetable](#) [Citation](#)<sup>new</sup> [Organizers](#) [Contact](#) [Workshop](#) [Download](#) [Evaluation Server](#)

### News


- September 2, 2014: [A new paper](#) which describes the collection of the ImageNet Large Scale Visual Recognition Challenge dataset, analyzes the results of the past five years of the challenge, and even compares current computer accuracy with human accuracy is now available. *Please cite it when reporting ILSVRC2012 results or using the dataset.*
- March 19, 2013: Check out [ILSVRC 2013!](#)
- January 26, 2012: [Evaluation server](#) is up. Now you can evaluate you own results against the competition entries.
- December 21, 2012: [Additional analysis of the ILSVRC dataset and competition results is released.](#)
- October 21, 2012: Slides from the workshop are being added to the [workshop schedule](#).
- October 13, 2012: [Full results](#) are released.
- October 8, 2012: Preliminary results have been released to the participants. Please join us at the [PASCAL VOC workshop](#) on October 12 at ECCV 2012. The workshop schedule for ILSVRC 2012 is [here](#)
- September 17, 2012: The submission deadline has been extended to September 30, 2012 (Sunday, 23:00 GMT). There will be no more extension.
- September 11, 2012: The [submission server](#) is up. You can submit your results now!
- July 10, 2012: Test images are released.
- June 16, 2012: The development kit, training and validation data released. Please [register](#) to obtain the download links.
- May 29, 2012: Registration page is up! Please [register](#)
- May 7, 2012: We are preparing to run the ImageNet Large Scale Visual Recognition Challenge 2012 (ILSVRC2012). **New task this year: fine-grained classification on 120 dog sub-classes!** Stay tuned!

### Workshop Schedule


- 15:30 - 16:00. Introduction and overview of results. Fei-Fei Li [ [slides](#) ]
- 16:00 - 16:25. Invited talk. OXFORD\_VGG team [ [slides](#) ] NB: **This is unpublished work. Please contact the authors if you plan to make use of any of the ideas presented**
- 16:25 - 16:40. Break
- 16:40 - 17:05. Invited Talk. ISI team [ [slides](#) ] NB: **This is unpublished work. Please contact the authors if you plan to make use of any of the ideas presented**
- 17:05 - 17:30. Invited Talk. SuperVision team [ [slides](#) ]
- 17:30 - 18:00. Discussion.

# Encoder


- Using pretrain models to extract feature vector from a given input image
  - Using pretrained ResNet-152
  - From Torchvision




LeNet-5




Convolution networks



AlexNet




This is getting complicated



VGG-19



Deep learning



ResNet-34



WTF?

ResNet-152



WTF! ! !

# Torchvision

- Pytorch official package consists of
  - popular datasets
  - model architectures
  - common image transformations for computer vision.

The screenshot shows the PyTorch torchvision documentation page. The top navigation bar includes links for 'Get Started', 'Features', 'Ecosystem', 'Blog', 'Tutorials', 'Docs' (which is highlighted with a red dot), 'Resources', and 'Github'. Below the navigation bar, the left sidebar contains a search bar labeled 'Search Docs' and a list of links under the heading 'Notes', including 'Autograd mechanics', 'Broadcasting semantics', 'CPU threading and TorchScript inference', 'CUDA semantics', 'Extending PyTorch', 'Frequently Asked Questions', 'Features for large-scale deployments', 'Multiprocessing best practices', 'Reproducibility', 'Serialization semantics', and 'Windows FAQ'. At the bottom of the sidebar are links for 'Community' and 'PyTorch Contribution Guide'. The main content area on the right is titled 'TORCHVISION' and contains a paragraph stating: 'The `torchvision` package consists of popular datasets, model architectures, and common image transformations for computer vision.' Below this is a section titled 'Package Reference' which lists the following modules: `torchvision.datasets`, `MNIST`, `Fashion-MNIST`, `KMNIST`, `EMNIST`, `QMNIST`, `FakeData`, `COCO`, `LSUN`, `ImageFolder`, `DatasetFolder`, and `ImageNet`.

# Pretrained ResNet-152

- Pretrained on the ILSVRC-2012-CLS

```
import torchvision.models as models  
resnet = models.resnet152(pretrained=True)
```

- Delete the last fc layer, use **NEW linear layer** to transform feature vector to have the same dimension as the input dimension of the LSTM network

```
13 # delete the last fc layer.  
14 modules = list(resnet.children())[:-1]  
15 self.resnet = nn.Sequential(*modules)  
16 self.linear = nn.Linear(resnet.fc.in_features, embed_size)  
17 self.bn = nn.BatchNorm1d(embed_size, momentum=0.01)
```

# Pretrained ResNet-152

```
8 class EncoderCNN(nn.Module):
9     def __init__(self, embed_size):
10         """Load the pretrained ResNet-152 and replace top fc layer."""
11         super(EncoderCNN, self).__init__()
12         resnet = models.resnet152(pretrained=True)
13         # delete the last fc layer.
14         modules = list(resnet.children())[:-1]
15         self.resnet = nn.Sequential(*modules)
16         self.linear = nn.Linear(resnet.fc.in_features, embed_size)
17         self.bn = nn.BatchNorm1d(embed_size, momentum=0.01)
18
19     def forward(self, images):
20         """Extract feature vectors from input images."""
21         with torch.no_grad():
22             features = self.resnet(images)
23             features = features.reshape(features.size(0), -1)
24             features = self.bn(self.linear(features))
25         return features
```



# Parameters Update

- In train.py Line 45 ~ 46
  - ResNet part parameters won't update

```
43 # Loss and optimizer
44 criterion = nn.CrossEntropyLoss()
45 params = list(decoder.parameters()) + list(encoder.linear.parameters()) + list(encoder.bn.parameters())
46 optimizer = torch.optim.Adam(params, lr=args.learning_rate)
```

# Decoder

- Noticed that you will use your model in model.py line 34
  - You can use nn.LSTM to check your environment is OK or not

```
--
27 class DecoderRNN(nn.Module):
28     def __init__(self, embed_size, hidden_size, vocab_size, num_layers, max_seq_length=20):
29         """Set the hyper-parameters and build the layers."""
30         super(DecoderRNN, self).__init__()
31         self.embed = nn.Embedding(vocab_size, embed_size)
32         # uncomment this line to use the default setting
33         #self.lstm = nn.LSTM(embed_size, hidden_size, num_layers, batch_first=True)
34         self.lstm = my_LSTM(embed_size, hidden_size, num_layers, batch_first=True)
35         self.linear = nn.Linear(hidden_size, vocab_size)
36         self.max_seg_length = max_seq_length
37
```

# LSTM Recall

- At professor slide “RecurrentNeuralNetworks.pdf”

- Memory state:  $\mathbf{s}^{(t)}$
- Input gate:  $\mathbf{g}^{(t)} = \sigma(\mathbf{U}^g \mathbf{x}^{(t)} + \mathbf{W}^g \mathbf{h}^{(t-1)})$
- Output gate:  $\mathbf{q}^{(t)} = \sigma(\mathbf{U}^o \mathbf{x}^{(t)} + \mathbf{W}^o \mathbf{h}^{(t-1)})$
- Forget gate:  $\mathbf{f}^{(t)} = \sigma(\mathbf{U}^f \mathbf{x}^{(t)} + \mathbf{W}^f \mathbf{h}^{(t-1)})$
- New content:  $\mathbf{a}^{(t)} = \mathbf{U} \mathbf{x}^{(t)} + \mathbf{W} \mathbf{h}^{(t-1)}$
- Memory update:  $\mathbf{s}^{(t)} = \mathbf{f}^{(t)} \odot \mathbf{s}^{(t-1)} + \mathbf{g}^{(t)} \odot \tanh(\mathbf{a}^{(t)})$
- Hidden unit update:  $\mathbf{h}^{(t)} = \mathbf{q}^{(t)} \odot \tanh(\mathbf{s}^{(t)})$
- Output unit update:  $\mathbf{o}^{(t)} = \mathbf{V} \mathbf{h}^{(t)}$

# Lstm Implement Hint

- You can use nn.Linear to build your lstm

```
self.fc_ho = nn.Linear(hidden_size, hidden_size, bias=if_bias)
```

- RNN Example

# Training

```
mtk11243@colglx0010:/proj/gpu_atp3/lab_test/LAB3
```

```
$ python3 train.py
```

```
[Debug] device cuda
```

```
Namespace(batch_size=128, caption_path='data/annotations/captions_train2014.json', crop_size=  
dels/', num_epochs=5, num_layers=1, num_workers=2, save_step=1000, vocab_path='data/vocab.pkl'  
loading annotations into memory...
```

```
Done (t=4.41s)
```

```
creating index...
```

```
index created!
```

```
Epoch [0/5], Step [0/3236], Loss: 9.2050, Perplexity: 9947.2122
```

# Testing

```
mtk11243@colglx0010:/proj/gpu_atp3/lab_test/LAB3
$ python3 sample.py --image='png/example.png'
<start> a group of giraffes standing in a field . <end>
mtk11243@colglx0010:/proj/gpu_atp3/lab_test/LAB3
$ python3 sample.py --image='png/ext.jpg'
<start> a group of motorcycles are parked on the street . <end>
mtk11243@colglx0010:/proj/gpu_atp3/lab_test/LAB3
$ python3 sample.py --image='png/201003151731430.jpg'
<start> a man riding a skateboard on top of a building . <end>
```

# Report Spec & Demo

- Introduction (5%)
- Explain how you implement LSTM (45%)
- Results – generating corresponding descriptions
  - A. example.png (10%)
  - B. ext.png (10%)
- Discussion (10%)
- Demo
  - Test your model on a given picture (10%)
  - Question (10%)

# Demo example

- Like “Results” in Report
- The demo testing will be uploaded before Demo



<start> a group of people riding bikes down a street . <end>



# Future Topic

- Pytorch extension
  - Define your own layer
- Dataloader